

lab7

Credit: materials adapted from Patrick Chester, with some examples taken from Ken Benoit's NYU Dept. of Politics short course Fall 2014

```
# Clear Global Environment
rm(list = ls())

# Setting WD
setwd("/Users/Lingyi/TAD/lab/Text-as-Data-Lab-Spr2018/W7_03_08_18/")

# Installing / Loading Libraries
# install.packages("tm")
# install.packages("NLP")
# install.packages("https://cran.r-project.org/bin/windows/contrib/3.4/prodlim_1.6.1.zip",
# repos = NULL, method = "libcurl")
#install.packages("RTextTools")

library(NLP)
library(tm)
library(RTextTools)
```

```
## Loading required package: SparseM
##
## Attaching package: 'SparseM'
## The following object is masked from 'package:base':
##
##      backsolve
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

1 Visualizing Bullying Data—Example from Pablo Barbera's Short Course on R, NYU 2016

```
# https://github.com/pablobarbera/data-science-workshop

df.tweets <- read.csv("bullying.csv", stringsAsFactors = F)

# Identify posts with and without bullying traces and create large documents
no_bullying <- paste(df.tweets$text[df.tweets$bullying_traces=="n"], collapse=" ")
yes_bullying <- paste(df.tweets$text[df.tweets$bullying_traces=="y"], collapse=" ")

# Create DTM and preprocess
groups <- VCorpus(VectorSource(c("No bullying" = no_bullying, "Yes bullying" = yes_bullying)))
groups <- tm_map(groups, content_transformer(tolower))
```

```

groups <- tm_map(groups, removePunctuation)
groups <- tm_map(groups, stripWhitespace)
bullying_dtm <- DocumentTermMatrix(groups)

# Label the two groups
bullying_dtm$dimnames$Docs = c("No bullying", "Yes bullying")


# Transpose matrix so that we can use it with comparison.cloud
bullying_tdm <- t(bullying_dtm)

# Compute TF-IDF transformation
bullying_tdm <- as.matrix(weightTfIdf(bullying_tdm))

# Display the two word clouds
comparison.cloud(bullying_tdm, max.words=100, colors=c("red", "blue"))

```

No bullying



Yes bullying

```

# Function is from the wordcloud package

```

2 Classification with SVM

```

# A) Linear - whole sample

# Let's train an SVM
df.tweets$type <- as.numeric(factor(df.tweets$bullying_traces))

# New package, better for SVM

```

```

##?create_matrix
bullying_dfm <- create_matrix(df.tweets$text,
                             language="english",
                             stemWords = FALSE,
                             weighting = weightTfIdf,
                             removePunctuation = FALSE
                             )

str(bullying_dfm)

## List of 6
## $ i      : int [1:45029] 1 1 1 1 1 1 1 1 1 1 ...
## $ j      : int [1:45029] 1381 1698 1841 2470 3586 4250 4272 5097 7251 7533 ...
## $ v      : Named num [1:45029] 1.1176 0.0969 1.1176 1.0267 0.9736 ...
## ..- attr(*, "names")= chr [1:45029] "1" "1" "1" "1" ...
## $ nrow    : int 5022
## $ ncol    : int 11824
## $ dimnames:List of 2
## ..$ Docs : chr [1:5022] "#Canada News: And the saga of official bullying #bilingualism continues."
## ..$ Terms: chr [1:11824] "" "" "" "" ...
## - attr(*, "class")= chr [1:2] "DocumentTermMatrix" "simple_triplet_matrix"
## - attr(*, "weighting")= chr [1:2] "term frequency - inverse document frequency (normalized)" "tf-idf"

# Make it all in-sample
##?create_container
container <- create_container(bullying_dfm,
                             t(df.tweets$type),
                             trainSize = 1:length(df.tweets$type),
                             virgin = FALSE
                             )

# Train the model
##?cross_validate
cv.svm <- cross_validate(container, nfold = 2, algorithm = 'SVM', kernel = 'linear')

## Fold 1 Out of Sample Accuracy = 0.9002997
## Fold 2 Out of Sample Accuracy = 0.8843024

## Comments:
# linear vs radial kernels, radial can overfit
# linear kernel is faster
# nfold is the number of times you have a different test set

# B) Linear - 90% Training data

training_break <- as.integer(0.9*nrow(df.tweets))

container <- create_container(bullying_dfm,
                             t(df.tweets$type),
                             trainSize = 1:training_break,
                             testSize = (training_break+1):nrow(df.tweets),
                             virgin = FALSE
                             )

# Let's train the model
cv.svm <- cross_validate(container,

```

```

        nfold = 4,
        algorithm = 'SVM',
        kernel = 'linear'
    )

## Fold 1 Out of Sample Accuracy = 0.7858268
## Fold 2 Out of Sample Accuracy = 0.780083
## Fold 3 Out of Sample Accuracy = 0.8103586
## Fold 4 Out of Sample Accuracy = 0.8003096

# Validate
cv.svm$meanAccuracy

## [1] 0.7941445

prop.table(table(df.tweets$type)) # baseline

##
##          1          2
## 0.7369574 0.2630426

# How well did we do?

# C) Radial - 90% training data

# Let's try again with the radial kernel
cv.svm <- cross_validate(container,
                        nfold = 4,
                        algorithm = 'SVM',
                        kernel = 'radial'
    )

## Fold 1 Out of Sample Accuracy = 0.730738
## Fold 2 Out of Sample Accuracy = 0.7667785
## Fold 3 Out of Sample Accuracy = 0.7333861
## Fold 4 Out of Sample Accuracy = 0.7419355

cv.svm$meanAccuracy

## [1] 0.7432095

# D) Linear - 50% training data

# What if we try with different % test/train?
training_break <- as.integer(0.5*nrow(df.tweets))

# There is no theoretical reason to choose .5 or .9
container <- create_container(bullying_dfm,
                             t(df.tweets$type),
                             trainSize = 1:training_break,
                             testSize = (training_break+1):nrow(df.tweets),
                             virgin = FALSE
    )
cv.svm$meanAccuracy

## [1] 0.7432095

```

```
prop.table(table(df.tweets$type)) # baseline
```

```
##  
##           1           2  
## 0.7369574 0.2630426
```

3 Virality of stories from NYT

```
nyt.fb <- read.csv("nyt-fb.csv", stringsAsFactors = FALSE)  
  
#str(nyt.fb)  
  
# Create variables for month and hour  
#head(nyt.fb$created_time)  
  
month <- substr(nyt.fb$created_time, 6, 7)  
  
hour <- substr(nyt.fb$created_time, 12, 13)  
  
nyt.fb <- data.frame(nyt.fb, month, hour)  
  
# Create a "viral" index  
total.resp <- nyt.fb$likes_count + nyt.fb$shares_count + nyt.fb$comments_count  
  
# Look at the extreme of the distribution  
perc_90 <- quantile(total.resp, .9)  
  
# Create a binary y variable with values 2 being viral and 1 being non-viral  
nyt.fb$viral <- as.numeric(total.resp > perc_90)  
  
# For the purposes of not destroying my laptop, let's choose a set of features  
training_break <- as.integer(0.9*nrow(nyt.fb))  
  
# A) Classification with SVM  
nyt_dtm <- create_matrix(nyt.fb$message, language="english", stemWords = FALSE,  
                        weighting = weightTfIdf, removePunctuation = FALSE)  
  
container <- create_container(nyt_dtm, t(nyt.fb$type), trainSize=1:training_break,  
                             testSize=(training_break+1):nrow(nyt.fb), virgin=FALSE)  
  
cv.svm <- cross_validate(container, nfold = 2, algorithm = 'SVM', kernel = 'linear')  
  
## Fold 1 Out of Sample Accuracy = 0.9536051  
## Fold 2 Out of Sample Accuracy = 0.9540253  
cv.svm$meanAccuracy  
  
## [1] 0.9538152  
prop.table(table(nyt.fb$viral))  
  
##
```

```
##      0      1
## 0.9 0.1
```

```
# B) Classification with logistic regression
```

```
message <- removePunctuation(tolower(nyt.fb$message))
nyt.fb$israel <- grepl("israel", message)
nyt.fb$trump <- grepl("trump", message)
nyt.fb$hillary <- grepl("hillary", message)
nyt.fb$obama <- grepl("barack|obama", message)
nyt.fb$terror <- grepl("terror|isis|isil|qaeda", message)
nyt.fb$kill <- grepl("kill|murder|shot", message)
nyt.fb$debate <- grepl("debat", message)
```

```
# Fitting a logistic model
```

```
glm.viral <- glm(as.factor(viral) ~ month + hour +
                 israel + trump + hillary + obama + terror + kill +
                 debate, data=nyt.fb, family=binomial(logit))
```

```
tab <- table(round(glm.viral$fitted.values), nyt.fb$viral)
```

```
# Accuracy of Logistic Regression
```

```
sum(diag(tab))/sum(tab)
```

```
## [1] 0.8998996
```

```
# In this case, SVM had a higher level of accuracy. Why might that be?
```